

# 400 bags of grocery receipts

+

# Neo4j

...

An exploration of Instacart data using Neo4j

# **3 Million Instacart Orders, Open Sourced**

Curious about the food Americans eat? Look no further.







Data and Data Modeling

Mechanics of getting data into Neo4j

Dumpster Diving

Production Concerns (Sizing, import strategy)

```
1 CREATE (i:Person {name: "Jonathan Freeman", twitter: "freethejazz"}),
2       (i)-[w:WORKS_FOR]->(c:Company {name: "Spantree Technology Group", twitter: "spantreeellc"}),
3       (i)-[l1:LIKE]->(a:Activity {name: "Cooking"}),
4       (i)-[l2:LIKE]->(a2:Activity {name: "Playing Music"})
5 RETURN i, c, a, a2, l1, l2
```



Graph



Table



Text



Code

\*(4)

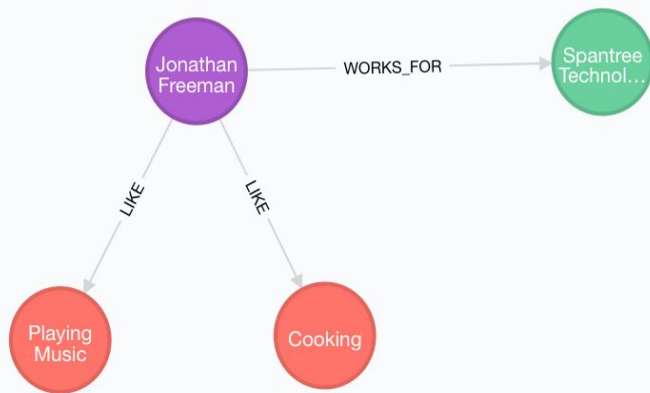
Person(1)

Company(1)

Activity(2)

\*(2)

LIKE(2)



# Data and Data Modeling

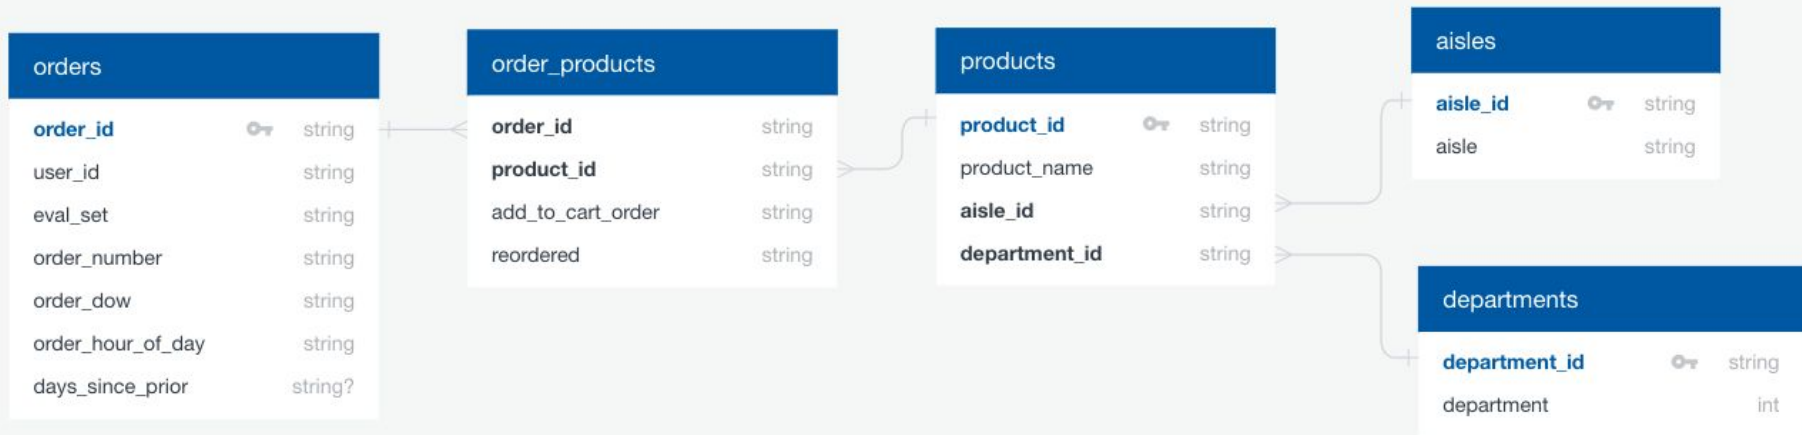
Mechanics of getting data into Neo4j

Dumpster Diving

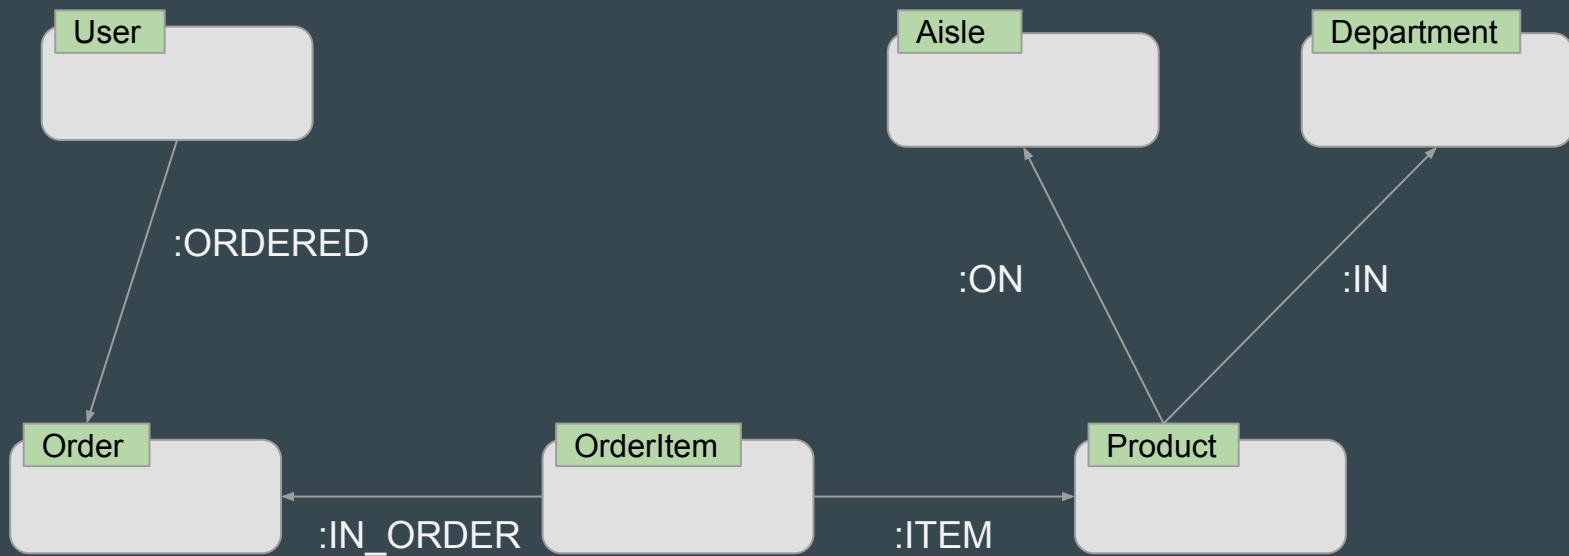
Production Concerns (Sizing, import strategy)



# CSV Model



# Graph Model



~~Data and Data Modeling~~

Mechanics of getting data into Neo4j

Dumpster Diving

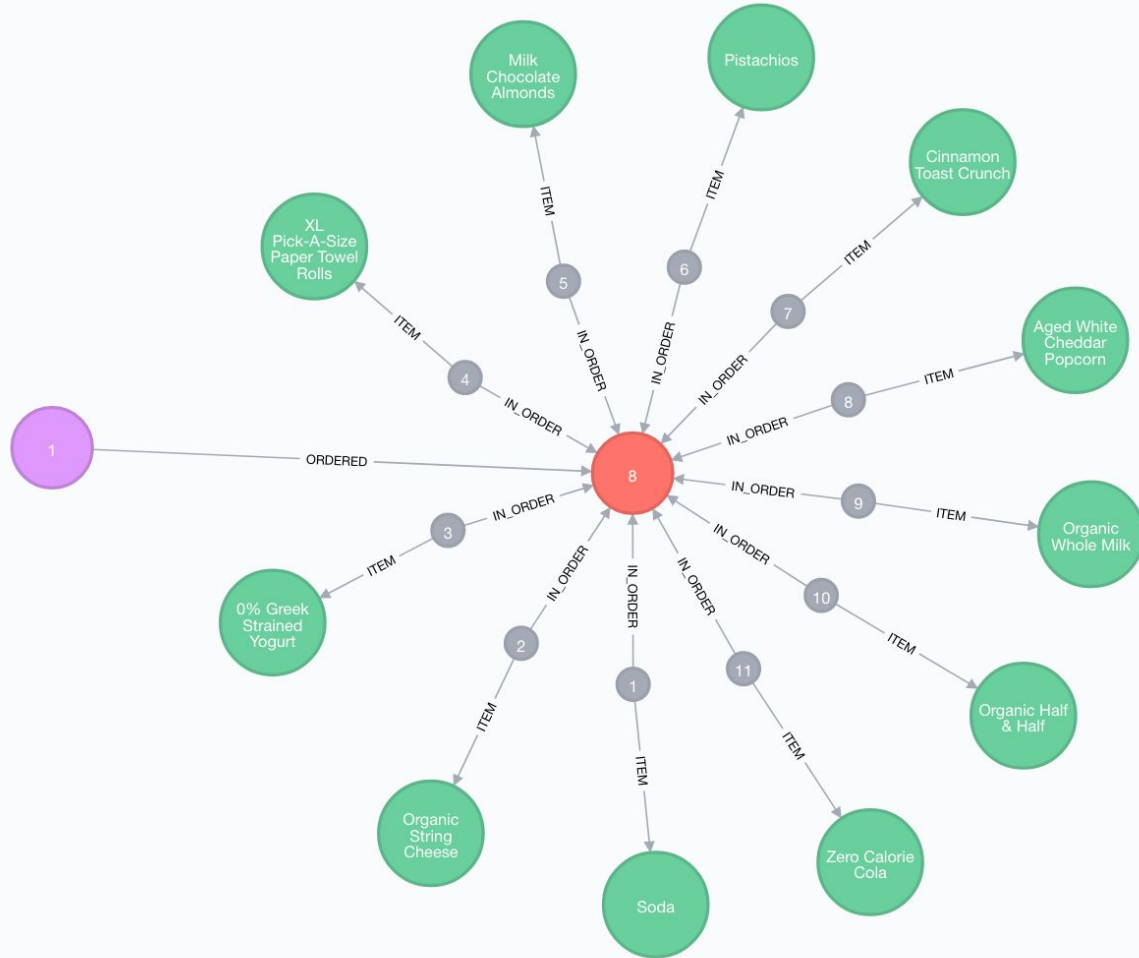
Production Concerns (Sizing, import strategy)


# IMPORT FROM CSV WITH HEADERS

```
1 LOAD CSV WITH HEADERS FROM "file:///departments.csv" as line
2 CREATE (:Department {id: line.department_id, name: line.department})
```

# Works with relationships, too

```
1 USING PERIODIC COMMIT
2 LOAD CSV WITH HEADERS FROM "file:///products.csv" as line with line
3 ⚠ match (a:Aisle {id: line.aisle_id}), (d:Department {id: line.department_id})
4 MERGE (a)<-[:ON]-(p:Product {id: line.product_id, name: line.product_name})-[:IN]->(d)
```





~~Data and Data Modeling~~

~~Mechanics of getting data into Neo4j~~

Dumpster Diving


Production Concerns (Sizing, import strategy)




# Dietary Restrictions

# Get Departments


\$ MATCH (d:Department) return d.name



Table



Text



Code

"d.name"
"frozen"
"other"
"bakery"
"produce"
"alcohol"
"international"
"beverages"
"pets"
"dry goods pasta"
"bulk"
"personal care"
"meat seafood"
"pantry"
"breakfast"
"canned goods"
"dairy eggs"

Started streaming 21 records after 1 ms and completed after 1 ms.

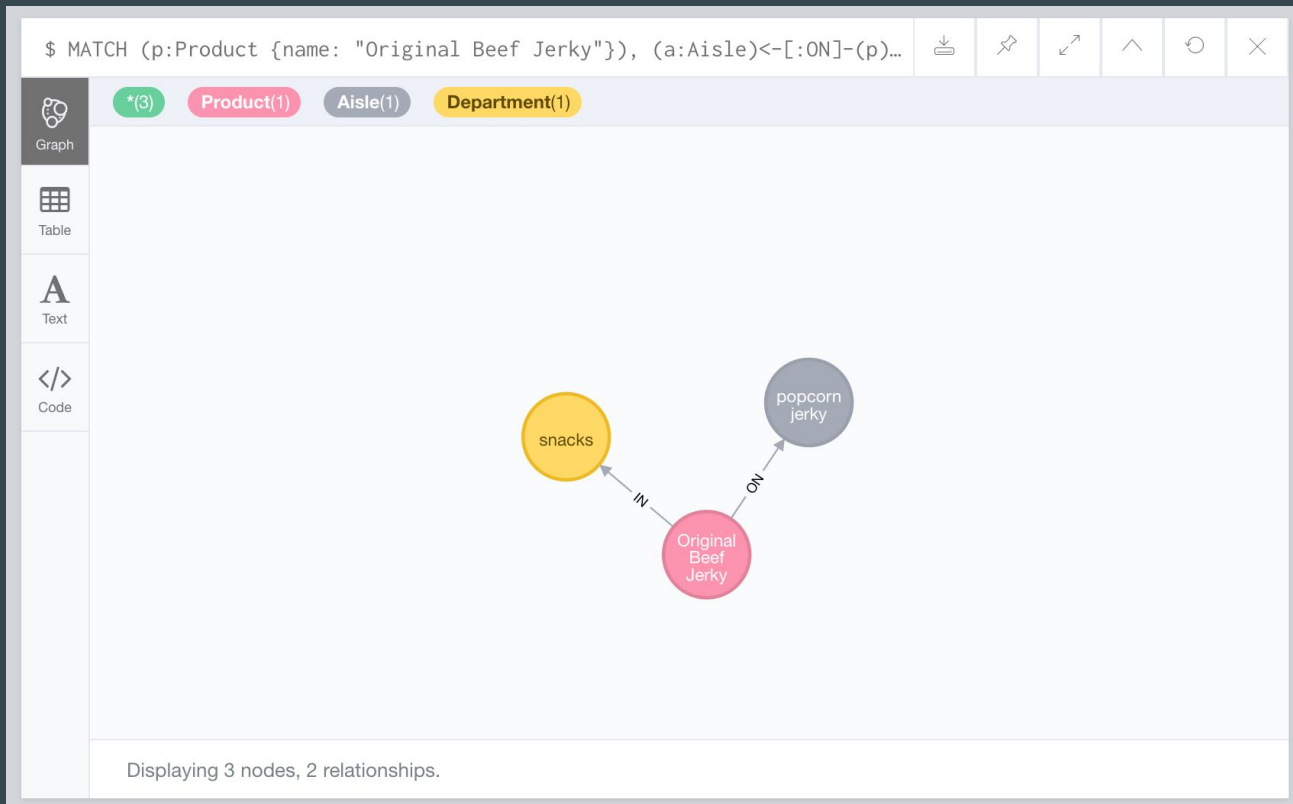
# Department-based Vegetarian

\$ MATCH (u:User) WITH u LIMIT 1 MATCH (u)-[:ORDERED]->(:Order)<-[:IN\_OR...

"p.name"	"overlap"
"Soda"	11
"Pistachios"	10
"Original Beef Jerky"	10
"Organic String Cheese"	9
"Zero Calorie Cola"	4
"Cinnamon Toast Crunch"	4
"Organic Half & Half"	3
"Aged White Cheddar Popcorn"	3
"XL Pick-A-Size Paper Towel Rolls"	3
"0% Greek Strained Yogurt"	2
"Bag of Organic Bananas"	2
"Milk Chocolate Almonds"	2
"Creamy Almond Butter"	1
"Bartlett Pears"	1
"Honeycrisp Apples"	1
"Organic Unsweetened Vanilla Almond Milk"	1

Started streaming 19 records after 15 ms and completed after 15 ms.

# Original Beef Jerky



# Aisles with Meat or Seafood in the name

```
$ MATCH (a:Aisle) WHERE a.name CONTAINS 'meat' OR a.name CONTAINS 'seaf...
```

Table

Text

Code

"a.name"
"marinades meat preparation"
"packaged meat"
"tofu meat alternatives"
"frozen meat seafood"
"canned meat seafood"
"lunch meat"
"meat counter"
"packaged seafood"
"seafood counter"

Started streaming 9 records after 3 ms and completed after 6 ms.

# Aisles to Avoid

\$ match (a:Aisle)<-[:ON]-([:Product]-[:IN]->(d:Department) where (a.name...

Table

Text

Code

"a.name"

"frozen meat seafood"

"packaged meat"

"packaged seafood"

"poultry counter"

"seafood counter"

"packaged poultry"

"hot dogs bacon sausage"

"meat counter"

"canned meat seafood"

"popcorn jerky"

"lunch meat"

Started streaming 11 records after 217 ms and completed after 360 ms.

\$ match (a:Aisle)<-[:ON]-(:Product)-[:IN]->(d:Department) where (a.name...



Table



Text



Code

"p.name"

"overlap"

"Banana"

240

"Bag of Organic Bananas"

209

"Organic Baby Spinach"

138

"Organic Strawberries"

123

"2% Reduced Fat Milk"

93

"Strawberries"

89

"Half & Half"

82

"Organic Hass Avocado"

82

"Organic Blueberries"

78

"Organic Whole Milk"

70

"Soda"

69

"Unsweetened Original Almond Breeze Almond Milk"

67

"Seedless Red Grapes"

67

"Hass Avocados"

66

"Organic Large Extra Fancy Fuji Apple"

66

"0% Greek Strained Yogurt"

66

Started streaming 3924 records after 1247 ms and completed after 1253 ms, displaying first 1000 rows.

# How Many Vegetarians?

```
$ match (a:Aisle)<-[:ON]-(:Product)-[:IN]->(d:Department) where (a.name...
```



Table

**"count(u)"**

52019



Text



Code

Started streaming 1 records after 268785 ms and completed after 268785 ms.



# Evolve the model

```
1 match (a:Aisle)<-[:ON]-(:Product)-[:IN]->(d:Department)
2 where (a.name contains 'meat'
3         or a.name contains 'seafood'
4         or d.name contains 'meat'
5         or a.name contains 'jerky')
6 and not a.name contains 'alternatives'
7 and not a.name contains 'marinades'
8 with collect(DISTINCT a.name) as avoidAisles
9 match (u:User)-[:ORDERED]->(:Order)<-[:IN_ORDER]-(:OrderedItem)-
    [:ITEM]->(:Product)-[:ON]->(a:Aisle)
10 WITH u, collect(DISTINCT a.name) as aisles, avoidAisles
11 WHERE none(a IN aisles WHERE a IN avoidAisles)
12 SET u:Vegetarian
```

# Add a label, et voila

```
$ MATCH (u:Vegetarian) return count(*)
```



Table

**"count(\*)"**

52019




Text



Code

Started streaming 1 records after 1 ms and completed after 1 ms.

# Products “Vegans” buy

\$ match (a:Aisle)-[:ON]-(:Product)-[:IN]->(d:Department) where (a.name...								
 Table	<b>"p.name"</b>	<b>"overlap"</b>						
	"Bag of Organic Bananas"	256						
	"Soda"	202						
	"Spring Water"	183						
	"Banana"	145						
	"Raspberries"	121						
	"Organic Strawberries"	113						
	"Clementines"	113						
	"Hass Avocados"	93						
	"Zero Calorie Cola"	89						
	"Organic Blueberries"	83						
	"Organic Baby Spinach"	76						
	"Smartwater"	67						
	"Sparkling Natural Mineral Water"	65						
	"Extra Fancy Unsalted Mixed Nuts"	64						
	"Crunchy Oats 'n Honey Granola Bars"	61						
	"Blackberries"	59						
Started streaming 2988 records after 9764 ms and completed after 9775 ms, displaying first 1000 rows.								

# How Many “Vegans”?

```
$ match (a:Aisle)<-[:ON]-(:Product)-[:IN]->(d:Department) where (a.name...
```



Table

**"COUNT(u)"**

9110



Text



Code

Started streaming 1 records after 279755 ms and completed after 279755 ms.

# Vegans like me

```
1 MATCH (u:Vegan {id: "92"}) with u
2 MATCH (u)--(o:Order)--(oi:OrderedItem)--(p:Product),
3     (other:Vegan)--(:Order)--(:OrderedItem)--(p)
4 return u, other, count(*) as overlap
5 order by overlap desc limit 5
```

\$ MATCH (u:Vegan {id: "92"}) with u MATCH (u)--(o:Order)--(oi:OrderedIt...



Graph



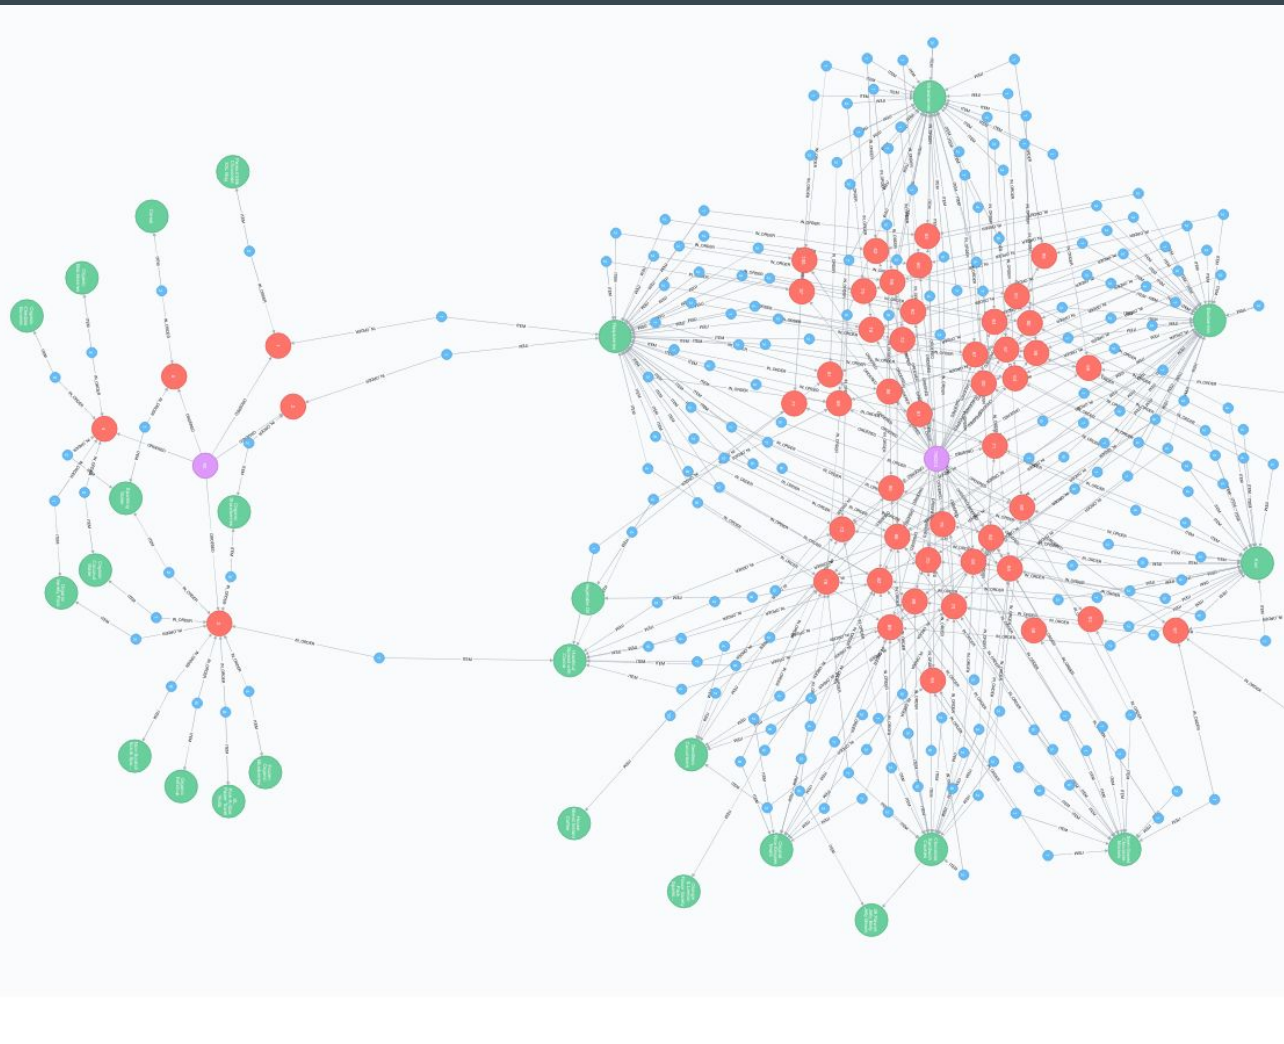
Table



Text



"u"	"other"	"overlap"
{"id": "92"}	{"id": "196224"}	190
{"id": "92"}	{"id": "168162"}	155
{"id": "92"}	{"id": "151588"}	131
{"id": "92"}	{"id": "41310"}	122
{"id": "92"}	{"id": "37866"}	120



# Vegans like me - Round 2

```
1 // Vegans Like me - Round 2
2 MATCH (u:Vegan {id: "92"}) with u
3 MATCH (u)--(o:Order)--(oi:OrderedItem)--(p:Product),
4       (other:Vegan)--(:Order)--(:OrderedItem)--(p)
5 return u, other, count(DISTINCT p) as overlap
6 order by overlap desc limit 5
```

\$ MATCH (u:Vegan {id: "92"}) with u MATCH (u)--(o:Order)--(oi:OrderedIt...



Graph



Table



Text



"u"

"other"

"overlap"

{"id": "92"}

{"id": "114638"}

6

{"id": "92"}

{"id": "86334"}

6

{"id": "92"}

{"id": "106828"}

5

{"id": "92"}


{"id": "121749"}

5

{"id": "92"}

{"id": "174108"}

5






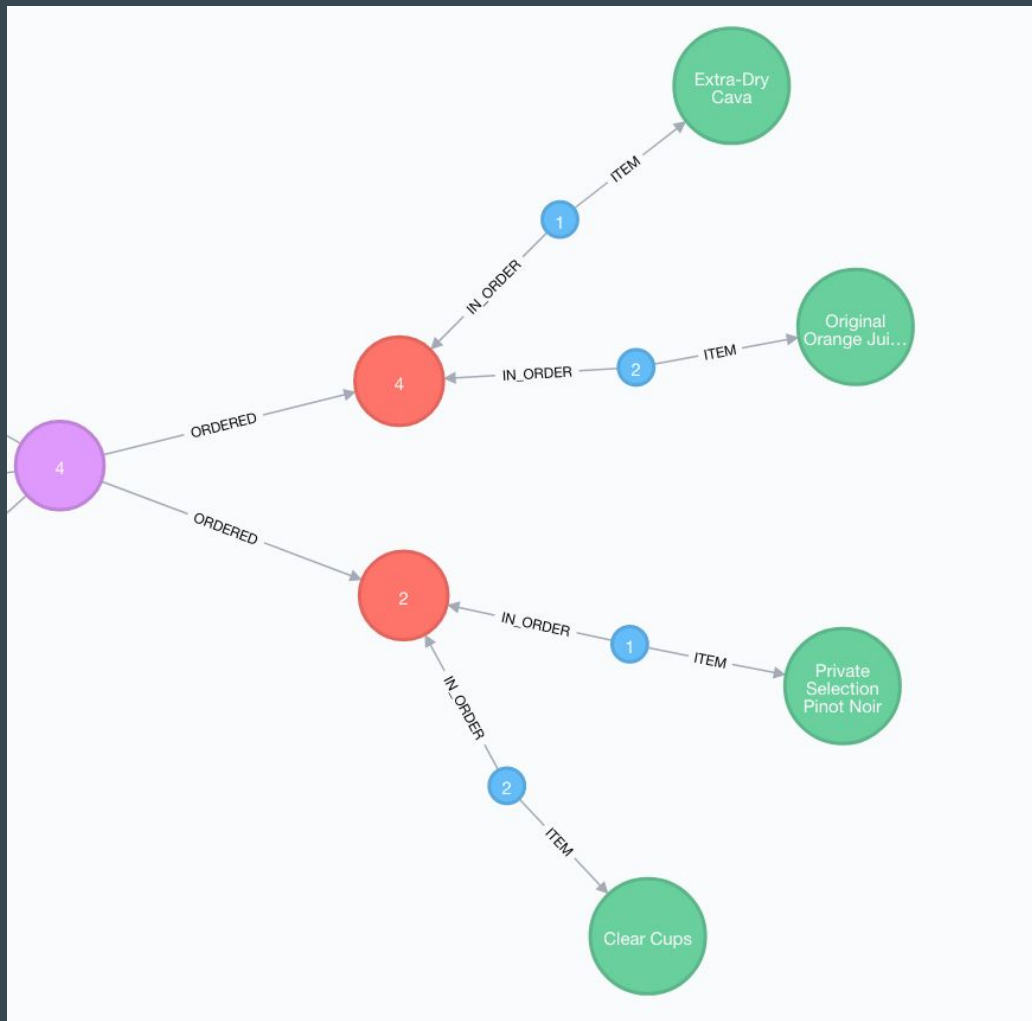


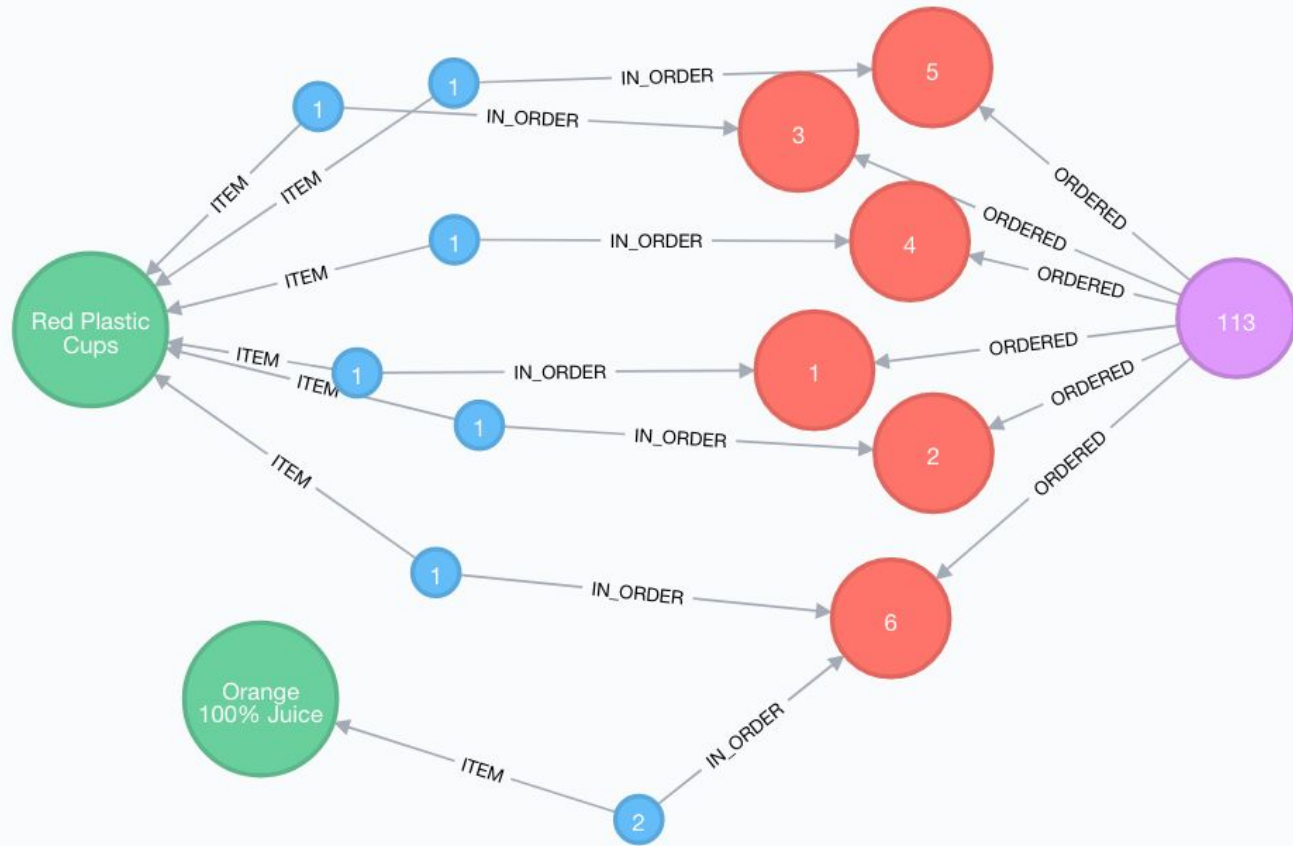
Instadate

One more thing...









# Cups and ....

\$ match (p:Product {name: "Red Plastic Cups"})<-[ITEM]-(oi:OrderedItem...

Table

Text

Code

"otherProduct.name"	"occurrences"
"Soda"	357
"Clementines"	189
"Bag of Organic Bananas"	173
"Zero Calorie Cola"	138
"Apples"	138
"Trail Mix"	126
"Reduced Fat 2% Milk"	124
"Mixed Fruit Fruit Snacks"	120
"Hass Avocados"	104
"Strawberries"	102
"Popcorn"	101
"Organic Simply Naked Pita Chips"	98
"Milk Chocolate Almonds"	91
"Crunchy Oats 'n Honey Granola Bars"	90
"Organic Tortilla Chips"	90
"0% Greek Strained Yogurt"	88

Started streaming 749 records after 44 ms and completed after 45 ms.

# Cups and Drinks

```
$ match (p:Product)<-[:ITEM]-(oi:OrderedItem)-[:IN_ORDER]->(o:Order), (...
```



Table



Text



Code

**"otherProduct.name"**

**"repeats"**

"Sauvignon Blanc"

58

"India Pale Ale"

54

"Beer"

43

"Cabernet Sauvignon"

40

"Chardonnay"

39

"Vodka"

30

"Prosecco Sparkling Wine"

26

"Pinot Noir"

26

"Chardonnay Wine"

25

"Pinot Grigio"

23

"Premium Lager Beer"

21

"Premium Belgian Lager"

19

"Red Blend"

19



# What else could you want?

\$ match (o:Order)<-[:IN\_ORDER]-(oi:OrderedItem)-[opr:ITEM]->(p:Product ...

Table

A

Text

</>

Code

"otherp.name"	"freq"
"Cabernet Sauvignon"	12
"India Pale Ale"	12
"Beer"	11
"Premium Belgian Lager"	8
"Belgian White Wheat Ale"	8
"Pinot Noir"	7
"Pinot Grigio"	7
"Chardonnay"	7
"Cabernet Sauvignon Wine"	5
"Belgian White Beer"	5

Started streaming 10 records after 38 ms and completed after 38 ms.

# Recommendations

# Problem w/ recommendations

```
$ match (o:Order)<-[:IN_ORDER]-(o:OrderItem)-[:ITEM]->(p:Product
```



Table



Text



Code

**"otherp.name"**

**"freq"**

"Clementines"

270

"Organic Simply Naked Pita Chips"

233

"Strawberries"

225

"Soda"

216

"Bag of Organic Bananas"

209

"Chocolate Chip Cookies"

204

"Cheez-It Cheddar Cracker"

204

"Trail Mix"

189

"Raspberries"


187

"Hass Avocados"

186

Started streaming 10 records after 84 ms and completed after 84 ms.

# Product Adjacency



```
1 // Build the product adjacency graph!
2 MATCH (u:Vegan)-[:ORDERED]->(o:Order)
3 WHERE o.evalSet <> "test"
4 with o
5 MATCH (o)<-[:IN_ORDER]-(oi:OrderedItem)-[:ITEM]->(p:Product)
6 with o, p
7 ORDER BY oi.addToCartOrder
8 with o, collect(p) as orderedProducts
9 UNWIND range(0, size(orderedProducts)-2) as i
10 WITH orderedProducts[i] as p1, orderedProducts[i+1] as p2
11 MERGE (p1)-[:NEXT]->(p2)
12   ON CREATE SET r.count = 1
13   ON MATCH SET r.count = r.count + 1
```

# What to buy after Chocolate Cookies?

```
$ match (p:Product {name: "Chocolate Sandwich Cookies"})-[r:n
```



Table



Text



Code

**"next.name"**

**"r.count"**

"Original Rice Krispies Treats"

15

"Blueberries"

7

"Soda"

5

"Pistachios"

4

"Raspberries"

4

"Organic Guacamole"

3

"Organic Tortilla Chips"

3

"Semi-Sweet Chocolate Morsels"

3

"Red Plastic Cups"

2

"Cheez-It Baked Snack Crackers"

2

Started streaming 10 records after 1 ms and completed after 1 ms.

# What people buy around Chocolate Cookies?

```
1 match (p:Product {name: "Chocolate Sandwich Cookies"})-[r:NEXT]->(next:Product)
2 return next.name, r.count
3 order by r.count desc
4 limit 10
```



```
1 match (p:Product {name: "Chocolate Sandwich Cookies"})-[r:NEXT]-(next:Product)
2 return next.name, sum(r.count) as count
3 order by count desc
4 limit 10
```

# What people buy around Chocolate Cookies?

```
$ match (p:Product {name: "Chocolate Sandwich Cookies"})-[r:N
```



Table



Text



Code

"next.name"

"count"

"Original Rice Krispies Treats"

21

"Semi-Sweet Chocolate Morsels"

18

"Red Plastic Cups"

9

"Soda"

8

"Blueberries"

7

"Chocolate Chip Cookies"

6

"Organic Guacamole"

6

"Strawberries"

5

"Milk Chocolate Almonds"

5

"Hazelnut Spread with Cocoa"

5

Started streaming 10 records after 1 ms and completed after 1 ms.



**What Else?**

What do people like to eat/drink  
when they're sick?

Match a person ordering frozen meals with  
someone ordering the ingredients in those  
meals

Can you find out if people have  
“unhealthy” habits? i.e., order flu  
medicine and a bottle of whiskey

Match people who just ordered pasta and  
sauce with someone who just ordered a  
bottle of red wine

Find out if it's a single person,  
a couple, or a family

~~Data and Data Modeling~~

~~Mechanics of getting data into Neo4j~~

~~Dumpster Diving~~

Production Concerns (Sizing, import strategy)

# Production Concerns

1. Hardware Sizing
2. Initial Data Load Strategy

# Make sure you have the memory!

Data size on disk for FS cache + 8-16 GB JVM heap + 1 GB for OS and misc stuff

# Total Memory Ballpark

1. Load a representative subset of the data
2. Extrapolate for full data size on disk

**34,000 Orders (~1%)**

Total Nodes: 412k

Total Rels: 1.5m

Total Props: 1.2m

Total Size on Disk: 333 MB

# Extrapolated full set

Total Nodes: 41m

Total Rels: 150m

Total Props: 120m

Total Size on Disk: 33 GB



# Total Memory Ballpark

Data size on disk for FS cache + 8 - 16 GB Java heap + 1 GB for OS and misc stuff

33 GB + 16 GB + 1GB = 50GB

m4.4xl with 64GB of memory for \$0.80/hour

## Extrapolated full set

Total Nodes: 41m -> 40.7m

Total Rels: 150m -> 71m

Total Props: 120m -> 128m

Total Size on Disk: 33 GB -> 5.18GB

# Importing a larger dataset

LOAD CSV

# Importing a larger dataset

~~LOAD CSV~~

neo4j-import

```
neo4j-import --into instacart.db --id-type string \  
  --skip-duplicate-nodes true \  
  --bad-tolerance 3300000 \  
  --nodes:User users_headers.csv,orders.csv \  
  --nodes:Department departments.csv \  
  --nodes:Aisle aisles.csv \  
  --nodes:Product products_headers.csv,products.csv \  
  --nodes:Order orders_headers.csv,orders.csv \  
  --nodes:OrderedItem order_products_headers.csv,order_products__prior_id.csv,order_products__train_id.csv \  
  --relationships:ON products_aisles_headers.csv,products.csv \  
  --relationships:IN products_departments_headers.csv,products.csv \  
  --relationships:ORDERED users_orders_headers.csv,orders.csv \  
  --relationships:IN_ORDER order_products_order_headers.csv,order_products__prior_id.csv,order_products__train_id.csv \  
  --relationships:ITEM order_products_products_headers.csv,order_products__prior_id.csv,order_products__train_id.csv
```

IMPORT DONE in 2m 57s 492ms.

Imported:

40711115 nodes

71158671 relationships

128719459 properties

Peak memory usage: 923.11 MB

**Thanks!**

Questions?



# Resources

Instagram Blog Post

<https://tech.instacart.com/3-million-instacart-orders-open-sourced-d40d29ead6f2>

Data Set

<https://www.instacart.com/datasets/grocery-shopping-2017>

Neo4j Hardware Sizing

<https://neo4j.com/news/video-hardware-sizing-for-neo4j/>

# Data Citation

“The Instacart Online Grocery Shopping Dataset 2017”, Accessed from <https://www.instacart.com/datasets/grocery-shopping-2017> on May 7th, 2017